

Generation latency across top k — pure top-k

Means of 50 prompt-level averages; 95% CIs are retained in the summary table

